

# Real-Time Image Classifier

## 1. Project Goal and Overview

The primary objective of this project was to develop a robust deep learning model capable of classifying images of **fruits and vegetables** and to deploy this model as a user-friendly, real-time web application using Streamlit.

- **Dataset:** Custom **Fruits\_Vegetables** dataset (containing **36 classes**).
- **Model Type:** Custom **Convolutional Neural Network (CNN)**.
- **Deployment:** Streamlit Dashboard for real-time inference.

## 2. Data Preparation and Preprocessing

The data pipeline was crucial for ensuring the model's accuracy and ability to generalize to new, unseen images.

### A. Data Splitting and Loading

The dataset was organized into three distinct directories to create the necessary splits: `train`, `validation`, and `test`.

- **Loading Utility:** `tf.keras.utils.image_dataset_from_directory` was used to load images directly from the folder structure.
- **Parameters:**
  - `image_size`: All images were resized to a consistent size (e.g., 180\*180 pixels).
  - `batch_size`: 32
  - **Labels:** Automatically inferred as integer labels (sparse encoding).

### B. Normalization and Rescaling

Normalization is applied to standardize pixel values, which is necessary for optimal model training.

- **Method:** A dedicated `layers.Rescaling(1./255)` layer was used.
- **Function:** This scales the raw pixel values (which range from 0 to 255) to a floating-point range of **[0, 1]**

### C. Image Augmentation

To prevent overfitting and increase the diversity of the training data, an aggressive augmentation pipeline was implemented **only on the training set**.

Augmentation Layer	Function
RandomFlip	Randomly flips images horizontally or vertically.
RandomRotation(0.2)	Randomly rotates images by up to ±36°

## 3. Model Development and Training

A custom CNN was designed and trained to handle the complexity of the 36-class image classification task.

### A. Algorithm Selection and Architecture

A custom CNN was chosen to extract hierarchical features from the image data.

- **Structure:** Sequential model consisting of multiple blocks of Conv2D→MaxPooling2
- **Regularization: Dropout layers** (e.g., 0.5) were included before the final dense layers to mitigate overfitting.
- **Output Layer:** The final layer has **36 units** with a **softmax** activation function, providing a probability distribution over the 36 classes.

### B. Compilation and Training

The model was compiled and trained using standard, optimized settings.

- **Optimizer:** **Adam** (adaptive learning rate optimization).
- **Loss Function:** **sparse\_categorical\_crossentropy** (appropriate for integer-encoded labels).
- **Early Stopping:** A callback was used to monitor the validation loss and stop training if the loss did not improve after a set number of epochs (patience), saving the best performing weights.

## 4. Evaluation and Prediction

### A. Model Evaluation

The final model was evaluated on the unseen **test set** to determine its real-world performance.

- **Primary Metric:** Accuracy (tracked throughout training and final evaluation).
- **Detailed Metrics:** Code was executed to generate a **Classification Report** (Precision, Recall, F1-Score) and a **Confusion Matrix**, providing a detailed view of class-specific performance.

### B. Real-Time Prediction

The notebook includes utility code demonstrating the prediction workflow for a single image, which forms the basis for the Streamlit application:

1. Load and display the image.
2. Preprocess the image (resizing, normalization).
3. Expand dimensions to create a batch size of 1 (`tf.expand_dims`).
4. Call `model.predict()`.
5. Apply `tf.nn.softmax` to the output to get class probabilities.
6. Use `tf.argmax` to select the highest-scoring class index.
7. Display the predicted class name and confidence score

## 5. Deployment (Streamlit)

The final step involves using the trained model (`image_classifier_model.keras`) within a Streamlit application (`app.py`) to create an interactive web interface for real-time prediction. This allows users to upload their own images and receive instant, visualized classification results. Below are some snapshots of the deliverables:




Figure 1: Image uploader




Figure 2: Classification Example